

Report Kernel Rebooters: Vision Division

Marco Carraro

`marco.carraro.23@studenti.unipd.it`

Luca Pellegrini

`luca.pellegrini.6@studenti.unipd.it`

Francesco Vezzani

`francesco.vezzani.1@studenti.unipd.it`

February 20, 2026

1 Introduction

The overall goal of this project is to build a complete computer vision pipeline that can classify flower images and support further analysis tasks, such as health-state estimation and performance evaluation. In practical terms, the system is designed around a clear workflow: images are loaded from a structured dataset, visual descriptors are extracted, test images are compared against training references, and final predictions are produced.

This report remains a shared team document, so section titles for all project components are kept even when implementation details are still pending. The current revision documents the HOG and BoW modules according to their latest simplified implementation: both modules now expose compact extractor interfaces focused only on feature extraction, vocabulary creation (for BoW), and descriptor matching.

2 Preprocessing

TODO

3 SIFT

3.1 Code Structure Decision

SIFT logic is organized into:

- `SIFTExtractor` (`include/sift.h`, `src/sift.cpp`) for low-level feature extraction, descriptor matching, and match filtering operations.
- `sift_processing` module (`include/sift_processing.h`, `src/sift_processing.cpp`) for dataset-level training and testing workflows.

This separation maintains architectural consistency with other modules: the extractor class handles algorithmic operations while the processing module orchestrates dataset traversal, descriptor aggregation, and prediction logic. As a result, `main.cpp` remains focused on high-level coordination, and the SIFT components can be independently tested, optimized, or replaced.

The `SIFTExtractor` interface exposes essential operations: `extract(...)`, `matchDescriptors(...)`, `filterMatches(...)`, and `matchAndFilter(...)`. Timing instrumentation is included to support performance analysis during development and evaluation phases.

3.2 Feature Extraction Choices

- Input images are in grayscale to focus on gradient information rather than color variation.
- SIFT parameters use OpenCV defaults with optional tuning: configurable number of features, octave layers, contrast threshold, edge threshold, and Gaussian sigma.
- The default configuration (`nfeatures=0`) extracts all detected keypoints, providing maximum detail for the training phase.

The decision to use grayscale processing aligns with SIFT's design philosophy: the algorithm was originally developed for intensity-based matching, and color information typically adds minimal discriminative value while increasing computational cost.

3.3 Training Pipeline

The training workflow follows a clear sequence:

1. Extract SIFT descriptors from all healthy and diseased (if `use_diseased` flag is set) training images.
2. Aggregate descriptors per class using vertical concatenation (`cv::vconcat`).
3. Store the combined descriptor matrices in a class-indexed map for efficient test-time retrieval.

This pipeline choice reflects a bag-of-features approach: all descriptors from a given class are pooled together without explicit spatial structure. The rationale is simplicity and robustness the classifier relies on the statistical distribution of local features rather than their precise geometric arrangement.

An important implementation detail is the use of `cv::vconcat` for descriptor combination, which is significantly faster than iterative concatenation and produces a single contiguous memory block per class.

3.4 Matching and Classification Strategy

For each test image:

1. Extract SIFT descriptors.
2. Match against each class's descriptor pool using FLANN-based matching.
3. Filter matches based on distance threshold: $d_{match} < threshold \times d_{min}$.
4. Classify as the class with the maximum number of good matches.

3.5 Matcher Configuration

Two matcher implementations were tested:

- `BFMatcher` with L2 norm: exhaustive search, guaranteed optimal matches, but slower for large descriptor sets (around 25 minutes for a classification).
- `FlannBasedMatcher`: approximate nearest neighbor search using KD-trees, significantly faster with same accuracy (around 5 minutes for a classification).

The FLANN matcher was selected as the default after benchmarking showed 5x speedup with negligible impact on classification accuracy for this dataset.

3.6 Robustness Choices

- Empty image check before extraction with early return and error message.
- Empty descriptor validation before matching operations.
- Graceful handling of images with zero detected keypoints (skipped in testing loop).
- Timing instrumentation for extraction and matching phases to identify performance bottlenecks.

3.7 Performance Considerations

The implementation includes several deliberate trade-offs:

- **Descriptor aggregation:** All training descriptors are stored in memory. For very large datasets (thousands of images per class), this could exceed available memory. A potential mitigation would be descriptor subsampling or clustering-based aggregation.
- **Exhaustive matching:** Each test descriptor is matched against all training descriptors for a given class. This scales as $O(N_{test} \times N_{train})$. FLANN provides partial relief.

These trade-offs were chosen to maintain code clarity and establish a functional baseline before introducing more complex optimization strategies.

3.8 Tested Configurations

Several parameters were evaluated during development:

- **Use of diseased descriptors:** Including diseased training samples improved classification accuracy.
- **Distance threshold:** Values between 1.5 and 2.5 were tested for the ratio test-inspired filtering, with 2.0 providing a good balance between precision and recall.
- **Matcher choice:** FLANN-based matching provided a significant speedup with no loss in accuracy compared to brute-force matching.
- **Nearest Neighbor Distance Ratio:** The original Lowe's ratio test was implemented but found to be too restrictive for this dataset, leading to very few matches. The distance-based thresholding approach provided a more flexible filtering mechanism that retained more valid matches while still reducing false positives.
- **Limits in train descriptor:** To mitigate the number of descriptors difference between classes, a maximum number of descriptors per class was tested (first 15000, then 30000). The sampling of those descriptors was done in multiple ways: random sampling, uniform sampling and k-means clustering. The best results were obtained with random sampling but this method was not consistent across runs. The other methods did not provide any improvement in accuracy and, especially k-means was very time consuming. At the end, there is no limit on the number of train descriptors.

The final configuration selected for the evaluation phase was FLANN-based matching with a distance threshold of 2.0 and no limit on the number of training descriptors, as this provided the best overall accuracy while maintaining reasonable processing time.

Matcher	Dataset	Thresh.	Train Descriptor	Time	Acc.	Daisy	Dand.	Rose	Sunf.	Tulip	NoFl.
BF	Full	2,5	NULL	26:15:996	20,31%	0,00%	41,67%	8,33%	16,67%	41,67%	0,00%
BF	Healthy	2,5	NULL	03:04:159	15,63%	8,33%	41,67%	0,00%	25,00%	8,33%	0,00%
FLANN	Full	2,5	NULL	06:54:950	20,31%	0,00%	41,67%	8,33%	16,67%	41,67%	0,00%
FLANN	Full	2,0	NULL	05:25:163	23,44%	8,33%	41,67%	16,67%	16,67%	41,67%	0,00%
FLANN	Full	2,5	15k (rnd sample)	04:17:885	20,31%	0,00%	33,33%	8,33%	25,00%	41,67%	0,00%
FLANN	Full	1,5	15k (rnd sample)	03:57:923	25,00%	16,67%	41,67%	8,33%	33,33%	33,33%	0,00%
FLANN	Full	1,5	NULL	05:29:062	18,75%	16,67%	33,33%	8,33%	25,00%	16,67%	0,00%
FLANN+NNDR	Full	0,8	NULL	05:33:180	17,19%	25,00%	16,67%	8,33%	16,67%	25,00%	0,00%
FLANN+NNDR	Full	0,7	NULL	05:23:491	23,44%	33,33%	33,33%	0,00%	16,67%	41,67%	0,00%
FLANN+NNDR	Full	0,7	15k (unif sample)	02:59:029	20,31%	41,67%	16,67%	8,33%	33,33%	8,33%	0,00%
FLANN+NNDR	Full	0,7	15k (k-means)	>30' train	18,75%	41,67%	25,00%	0,00%	25,00%	8,33%	0,00%
FLANN	Full	2,0	30k (unif sample)	04:02:434	20,31%	8,33%	41,67%	8,33%	25,00%	25,00%	0,00%

Table 1: SIFT performance evaluation

4 SURF

4.1 Code Structure Decision

SURF logic follows the same architectural pattern established for SIFT:

- `SURFExtractor` (`include/surf.h`, `src/surf.cpp`) for low-level feature extraction, descriptor matching, and match filtering operations.
- `surf_processing` module (`include/surf_processing.h`, `src/surf_processing.cpp`) for dataset-level training and testing workflows.

This parallel structure serves multiple purposes: it maintains consistency across the codebase, allows direct code comparison between SIFT and SURF implementations in `main.cpp`.

The `SURFExtractor` interface mirrors `SIFTExtractor`, exposing: `extract(...)`, `matchDescriptors(...)`, `filterMatches(...)`, and `matchAndFilter(...)`. This symmetry was a deliberate design decision to reduce cognitive load when working with multiple feature extractors and to facilitate comparative analysis.

4.2 Conditional Compilation

A critical implementation detail is that SURF support is optional and controlled through the `ENABLE_SURF` preprocessor flag. This is necessary because SURF is part of OpenCV’s `xfeatures2d` module, which requires separate installation and is not included in OpenCV’s default distribution due to patent restrictions.

The conditional compilation pattern is implemented using:

```
#ifdef ENABLE_SURF
// SURF implementation
#endif // ENABLE_SURF
```

This approach allows the project to compile and run without SURF when `xfeatures2d` is unavailable, while still supporting it when the module is properly installed. Users can enable SURF by configuring CMake with `-DCONFIG_ENABLE_SURF=ON`.

4.3 Feature Extraction Choices

- Input images are used in grayscale, consistent with SURF’s original design for intensity-based features.
- SURF parameters use sensible defaults: Hessian threshold 400.0, 4 octaves, 3 octave layers.

- The `extended` flag (defaulting to `false`) controls descriptor dimensionality: 64 dimensions for standard SURF, 128 for extended SURF.
- The `upright` flag (defaulting to `false`) determines rotation invariance: when false, descriptors are rotation-invariant but computationally more expensive.

The Hessian threshold is the primary parameter controlling the number of detected keypoints. Higher values detect fewer but more distinctive keypoints, while lower values increase keypoint count at the cost of potentially including less stable features. The default is set to 100.0 but it has been increased to 400.0 to reduce the number of keypoints, speed up processing and increasing accuracy.

4.4 Training Pipeline

The training workflow is architecturally identical to SIFT:

1. Extract SURF descriptors from healthy training images.
2. Optionally extract descriptors from diseased training images.
3. Aggregate descriptors per class using `cv::vconcat`.
4. Store combined descriptor matrices indexed by flower class.

This structural equivalence was intentional, it ensures that performance differences between SIFT and SURF can be attributed solely to the feature extraction algorithm rather than implementation variations in the classification pipeline.

4.5 Matching Strategy

The matching and classification strategy is identical to SIFT:

1. Extract SURF descriptors from test image.
2. Match against each class's descriptor pool using FLANN-based matching.
3. Filter matches: $d_{match} < \tau \cdot d_{min}$.
4. Classify as the class with maximum good matches.

Using the same threshold parameter τ across SIFT and SURF allows for direct comparison, though optimal threshold values may differ between the two algorithms due to differences in descriptor distance distributions.

5 ORB

5.1 Code Structure Decision

ORB follows the established architectural pattern:

- `ORBExtractor` (`include/orb.h`, `src/orb.cpp`) for low-level feature extraction, descriptor matching, and match filtering operations.
- `orb_processing` module (`include/orb_processing.h`, `src/orb_processing.cpp`) for dataset-level training and testing workflows.

This structural consistency was intentional: by maintaining identical interfaces and workflows across SIFT, SURF, and ORB modules, the codebase remains maintainable and direct algorithm comparisons become straightforward. The only algorithmic differences are encapsulated within the respective extractor classes.

5.2 Feature Extraction Choices

- Input images are used in grayscale, consistent with ORB's design for binary pattern matching.
- ORB parameters are set to default values apart from `nfeatures`.
- `nfeatures` is set to 1500, with threshold values adjusted accordingly for optimal classification performance.

A critical discovery during implementation was that ORB's `nfeatures` parameter behaves fundamentally differently from SIFT's equivalent. Setting `nfeatures=0` results in no keypoint detection, whereas SIFT interprets this as "extract all detected keypoints." This behavior necessitated explicit parameter specification and careful tuning to achieve acceptable performance.

5.3 Training Pipeline

The training workflow mirrors SIFT and SURF:

1. Extract ORB descriptors from healthy training images.
2. Optionally extract descriptors from diseased training images.
3. Aggregate descriptors per class using vertical concatenation.
4. Store combined descriptor matrices indexed by flower class.

This structural equivalence ensures that performance differences can be attributed solely to the feature extraction algorithm rather than implementation variations.

5.4 Matching Strategy

ORB matching follows the same high-level strategy as SIFT/SURF but with a critical difference in the underlying matcher:

1. Extract ORB descriptors from test image.
2. Match against each class's descriptor pool using BFMatcher with Hamming distance. FLANN-based matching is possible using LSH (Locality Sensitive Hashing) but typically slower than BFMatcher for binary descriptors.
3. Filter matches: $d_{match} < \tau \cdot d_{min}$.
4. Classify as the class with maximum good matches.

The classification decision remains:

$$\hat{y}(x) = \arg \max_c |\text{good_matches}(x, c)|$$

5.5 Parameter Tuning Process

Initial implementation with `nfeatures=0` resulted in zero keypoint detection across all images. This revealed a fundamental difference between ORB and SIFT parameter interpretation:

- SIFT: `nfeatures=0` → extract all detected keypoints
- ORB: `nfeatures=0` → extract exactly zero keypoints
- `nfeatures ≥ 1000` to ensure sufficient descriptor density per image

6 Template Matching

TODO

7 HOG

7.1 Objective

The HOG branch was introduced to add a global shape-and-gradient based descriptor that is easy to interpret and relatively stable in many real-world image conditions. The idea is simple: each test image is converted into one feature vector, then compared with feature vectors extracted from training images. The predicted class is taken from the nearest training sample.

From an engineering perspective, the objective was not to create a heavily optimized model, but to produce a transparent baseline that is easy to debug, easy to validate, and fully consistent with the project architecture.

7.2 Code Structure Decision

HOG logic is split into:

- `HOGExtractor` (`include/hog.h`, `src/hog.cpp`) for low-level feature extraction and distance computation.
- `hog(...)` wrapper (`src/matching.cpp`) for dataset-level loop and prediction print.

This separation is important because it keeps responsibilities clean: `HOGExtractor` handles only algorithmic operations, while `hog(...)` handles dataset traversal and prediction flow. As a result, `main.cpp` remains focused on orchestration, and the module can be reused or replaced with minimal impact on the rest of the system.

In the current version, `HOGExtractor` was simplified to a minimal API: `extract(...)` returns `bool` and `matchDescriptors(...)` returns the L2 distance. Timing fields and related getters were removed because they were not used by the application flow.

7.3 Feature Extraction Choices

- Input image is converted to grayscale if needed.
- The image is resized to a fixed window (64x128).
- HOG parameters are basic OpenCV defaults: block size 16x16, block stride 8x8, cell size 8x8, 9 bins.

The fixed resize step is a key practical decision: HOG vectors can only be compared directly when they share the same dimensionality. Using a fixed window gives deterministic descriptor length and avoids shape-dependent edge cases in matching.

7.4 Matching Choices

- For each test descriptor, all train descriptors are scanned.
- Similarity score is Euclidean distance (L2).
- Predicted label is the train sample with minimum distance.

This results in a straightforward nearest-neighbor baseline:

$$\hat{y}(x) = \arg \min_i \|h(x) - h(x_i^{train})\|_2$$

where $h(\cdot)$ is the HOG descriptor.

The benefit of this choice is interpretability: every prediction can be traced back to one specific training sample and one explicit distance value, which is very useful during qualitative inspection.

7.5 Robustness Choices

- Empty input image check in `extract(...)`.
- Empty or incompatible descriptor check before matching.
- Sample is skipped by the wrapper when `extract(...)` returns `false`.

7.6 Trade-off

The implementation is intentionally simple and readable, but the exhaustive comparison step has cost $O(N_{test} \cdot N_{train})$. This is acceptable as a baseline and for medium dataset sizes, but it is a known scalability limitation for larger datasets.

8 BoW

8.1 Objective

While HOG focuses on global gradient structure, BoW was introduced to capture local visual patterns through keypoints. The goal is to convert a variable number of local ORB descriptors into a fixed-size global representation, so images can be compared in a compact and uniform way.

In other words, BoW bridges local detail and global matching: it keeps the discriminative power of local descriptors but produces one standardized vector per image.

8.2 Code Structure Decision

BoW logic is split into:

- `BoWExtractor` (`include/bow.h`, `src/bow.cpp`) for vocabulary building, histogram extraction, and histogram distance.
- `bow(...)` wrapper (`src/matching.cpp`) for train/test loop and prediction print.

This separation follows the same architectural rule used across the project: core algorithm in an extractor class, orchestration in a lightweight wrapper. This improves readability and allows future experiments (different vocabularies, different local features, different distance metrics) without touching application flow.

The current version also simplifies the BoW extractor interface to three public operations: `buildVocabulary(...)`, `extract(...)`, and `matchDescriptors(...)`. Unused timing/keypoint getter APIs were removed to keep the code shorter and easier to maintain.

8.3 Pipeline Choices

The implemented BoW pipeline is:

1. extract ORB descriptors from all train images,
2. convert descriptors to CV_32F,
3. run k-means to build a visual vocabulary (default size: 20 words),
4. assign each descriptor to the closest visual word,
5. build one histogram per image and normalize it,
6. compare test and train histograms with L2 distance.

This pipeline is intentionally classical and explicit. Every stage can be inspected independently (descriptors, vocabulary, histogram), making debugging and incremental improvement easier for the team.

8.4 Why ORB + K-means

- ORB is already present in the project and is computationally light.
- K-means gives a direct and standard way to form visual words.
- Fixed vocabulary keeps the implementation deterministic and easy to tune.

Choosing ORB also reduces integration friction, since ORB-based utilities were already available and familiar in the codebase. K-means, despite being simple, provides a clear semantic interpretation: each cluster center is treated as a visual word.

8.5 Histogram and Matching Formula

Each image is represented by a normalized histogram $b(x) \in R^K$, where K is vocabulary size. Prediction follows:

$$\hat{y}(x) = \arg \min_i \|b(x) - b(x_i^{train})\|_2$$

Normalization reduces sensitivity to raw keypoint count differences between images, so comparison focuses more on visual word distribution than on absolute descriptor count. This is especially useful when images produce very different numbers of detected keypoints.

8.6 Robustness Choices

- Vocabulary availability check before extraction.
- Empty descriptor/histogram checks.
- Skip logic in `bow(...)` when `extract(...)` returns `false`.

8.7 Trade-off

The current implementation favors clarity over optimization:

- vocabulary is rebuilt at each run,
- nearest-neighbor matching is exhaustive,
- no TF-IDF or advanced scoring yet.

These trade-offs are deliberate at this stage: the code remains concise and easy to reason about, and it provides a reliable baseline before introducing acceleration or more advanced weighting schemes.

9 Metrics implementation

9.1 Code Structure Decision

The metrics system is organized into three components:

- `metrics.h/cpp` (`include/metrics.h`, `src/metrics.cpp`) for core data structures and computation functions.
- `print_stats.h/cpp` (`include/print_stats.h`, `src/print_stats.cpp`) for formatted output and reporting.
- `Metrics` struct as the central data container, passed by reference throughout the pipeline.

This separation permits modular development: `metrics.cpp` handles numerical computations and data management, while `print_stats.cpp` focuses exclusively on presentation formatting. This design makes it easy to add new output formats without modifying the core metrics logic.

9.2 Core Data Structure

The `Metrics` struct encapsulates all evaluation data:

```
struct Metrics {  
    int num_classes;  
    int total_samples;  
    int correct_predictions;  
    std::vector<std::vector<int>> confusion_matrix;  
    std::vector<double> processing_times;  
};
```

Key design decisions:

- **Confusion matrix:** Stored as `vector<vector<int>>` rather than a flat array for intuitive indexing: `confusion_matrix[true_class][predicted_class]`.
- **Processing times:** Stored as individual measurements rather than pre-computed aggregates, allowing flexible statistical analysis (mean, median, percentiles) without data loss.
- **Counts vs rates:** Raw counts (`correct_predictions`, `total_samples`) are stored; accuracy is computed on-demand to avoid synchronization issues during incremental updates.

9.3 Accuracy Computation

Multiple accuracy metrics are supported:

Overall accuracy:

$$\text{accuracy} = \frac{\text{correct_predictions}}{\text{total_samples}}$$

Per-class accuracy:

$$\text{accuracy}_c = \frac{\text{confusion_matrix}[c][c]}{\sum_j \text{confusion_matrix}[c][j]}$$

Per-class accuracy is computed by dividing the diagonal element (correct predictions for class c) by the sum of the entire row (all predictions where the true class was c). This metric reveals which classes are well-distinguished and which are frequently confused.

The implementation includes proper handling of edge cases: zero-sample classes return 0.0 accuracy rather than causing division by zero.

9.4 Timing Statistics

Processing time analysis provides four key metrics:

- **Total time:** Sum of all per-image processing times.
- **Mean time:** Average processing time per image.
- **Min/Max time:** Identifies best-case and worst-case performance.

These metrics are computed as:

$$\text{mean} = \frac{1}{n} \sum_{i=1}^n t_i$$
$$\min = \min_i t_i, \quad \max = \max_i t_i$$

9.5 Reporting Functions

The `print_stats` module provides hierarchical reporting:

Component functions:

- `printConfusionMatrix(...)`: Formatted confusion matrix table.
- `printTimingStats(...)`: Processing time summary.
- `printPerClassAccuracy(...)`: Per-class accuracy breakdown.

Composite function:

- `printClassificationReport(...)`: Combines all three components into a comprehensive summary with optional algorithm name header.

This modular structure allows flexible reporting: debugging sessions might print only timing stats, while final evaluation uses the full classification report.

9.6 Design Patterns

Several design patterns enhance usability:

Const-correctness: Query functions take `const Metrics&` to communicate that they perform read-only operations:

```
double totalAccuracy(const Metrics& metrics);
```

Reference parameters: Update functions take `Metrics&` to enable efficient in-place modification:

```
void addPrediction(Metrics& metrics, int true_class, int predicted_class);
```

9.7 Extensibility

The metrics system is designed for easy extension:

- Adding new metrics: Extend the `Metrics` struct and add corresponding computation functions.
- New output formats: Add new functions to `print_stats.cpp` or create separate modules.
- Statistical tests: Raw data in `processing_times` vector enables computation of variance, standard deviation, percentiles, or hypothesis tests.

9.8 Performance Considerations

The metrics system adds minimal overhead:

- Per-prediction update: $O(1)$ for counter increments and confusion matrix access.
- Time recording: $O(1)$ vector push_back operation.
- Accuracy computation: $O(c^2)$ where c is number of classes, negligible compared to feature extraction time.

9.9 Integration with Feature Extractors

All feature extraction modules (SIFT, SURF, TM, HOG, BoW) use this common metrics infrastructure, ensuring:

- Consistent evaluation methodology across algorithms.
- Comparable timing measurements (all include feature extraction through final prediction).
- Uniform reporting format for easy comparison.

This standardization is critical for fair performance comparisons and for making informed decisions about which algorithm is best suited for the flower classification task.

10 Flower Health Detection

Section reserved for flower health detection details (to be completed).

11 Performance Evaluation

Section reserved for metrics, confusion matrix, and timing analysis (to be completed).

12 Conclusion

In this iteration, HOG and BoW were integrated as simple and modular baselines. The main design decision was to keep a consistent extractor interface and move algorithm-level loops outside `main.cpp` into a dedicated matching module. The latest revision further reduced complexity by removing unused state from HOG/BoW extractors and retaining only the operations actually used by the pipeline.

13 Work Distribution

The project workload was distributed among team members in the following way:

- Marco Carraro: SIFT, SURF, ORB and metrics system.
- Luca Pellegrini: preprocessing pipeline image container infrastructure, template matching.
- Francesco Vezzani: HOG and BoW implementation, matching wrappers,

All team members participated in code reviews, integration testing, and collaborative debugging sessions to ensure consistency across modules and resolve interface compatibility issues.

Hours worked	Marco Carraro	Luca Pellegrini	Francesco Vezzani
Total	~ 40	XX	XX

Table 2: Approximate hours worked per team member

References / Links

- Project repository: https://github.com/kekko7072/Final_Project_Kernel_Rebooters/tree/main
- OpenCV SIFT documentation: https://docs.opencv.org/4.5.4/d7/d60/classcv_1_1SIFT.html#a94ee0141f77675822e574bbd2c079811
- OpenCV SURF documentation: https://docs.opencv.org/4.6.0/d5/df7/classcv_1_1xfeatures2d_1_1SURF.html
- OpenCV FlannBasedMatcher documentation: https://docs.opencv.org/4.5.4/dc/de2/classcv_1_1FlannBasedMatcher.html
- OpenCV HOGDescriptor documentation: https://docs.opencv.org/4.x/d5/d33/structcv_1_1HOGDescriptor.html
- OpenCV ORB documentation: https://docs.opencv.org/4.x/db/d95/classcv_1_1ORB.html
- OpenCV kmeans documentation: https://docs.opencv.org/4.x/d1/d5c/tutorial_py_kmeans_opencv.html